

Algorithm 966: A Practical Iterative Algorithm for the Art Gallery Problem Using Integer Linear Programming

DAVI C. TOZONI, PEDRO J. DE REZENDE, and CID C. DE SOUZA, University of Campinas

In the last few decades, the search for exact algorithms for known NP-hard geometric problems has intensified. Many of these solutions use Integer Linear Programming (ILP) modeling and rely on state-of-the-art solvers to be able to find optimal solutions for large instances in a matter of minutes. In this work, we discuss an ILP-based algorithm that solves to optimality the Art Gallery Problem (AGP), one of the most studied problems in computational geometry. The basic idea of our method is to iteratively generate upper and lower bounds for the problem through the resolution of discretized versions of the AGP, which are reduced to instances of the Set Cover Problem. Our algorithm was implemented and tested on almost 3,000 instances and attained optimal solutions for the vast majority of them, greatly increasing the set of instances for which exact solutions are known. To the best of our knowledge, in spite of the extensive study of the AGP in the last four decades, no other algorithm has shown the ability to solve the AGP as effectively and efficiently as the one described here. Evidence of its robustness is presented through tests done on a number of classes of polygons of various sizes with and without holes. A software package implementing the algorithm is made available.

Categories and Subject Descriptors: I.3.5 [Computer Graphics]: Computational Geometry and Object Modeling—*Geometric Algorithms*; G.4 [Mathematical Software]: Algorithm design and analysis; G.1.6 [Mathematics of Computing]: Optimization—*Integer Programming*

General Terms: Design, Algorithms, Performance

Additional Key Words and Phrases: Art gallery problem, exact algorithm, combinatorial optimization, computational geometry

ACM Reference Format:

Davi C. Tozoni, Pedro J. de Rezende, and Cid C. de Souza. 2016. Algorithm 966: A practical iterative algorithm for the art gallery problem using integer linear programming. *ACM Trans. Math. Softw.* 43, 2, Article 16 (August 2016), 27 pages.

DOI: <http://dx.doi.org/10.1145/2890491>

1. INTRODUCTION

In 1973, Victor Klee proposed a new geometric puzzle, which basically consisted of answering how few guards would be sufficient to ensure that an art gallery, represented by a simple (i.e., non-self-intersecting) polygon, was fully guarded, assuming that a guard's field of view covers 360 degrees as well as an unbounded distance. This question became known as the *Art Gallery Problem* (AGP) and, in the course of time,

An outline of this algorithm was presented in a short proceedings paper of the 12th International Symposium on Experimental Algorithms (see Tozoni et al. [2013]).

This research was supported by grants from Conselho Nacional de Desenvolvimento Científico e Tecnológico (CNPq) #311140/2014-9, #477692/2012-5, and #302804/2010-2; Fundação de Amparo à Pesquisa do Estado de São Paulo (FAPESP) #2007/52015-0 and #2012/18384-7; and FAEPEX/UNICAMP.

Authors' addresses: D. C. Tozoni, P. J. de Rezende, and C. C. de Souza, Institute of Computing, University of Campinas, Av. Albert Einstein 1251, Cidade Universitária Zeferino Vaz, 13083-852, Campinas(SP), Brazil; emails: davi.tozoni@gmail.com, rezende@ic.unicamp.br, cid@ic.unicamp.br.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or permissions@acm.org.

© 2016 ACM 0098-3500/2016/08-ART16 \$15.00

DOI: <http://dx.doi.org/10.1145/2890491>

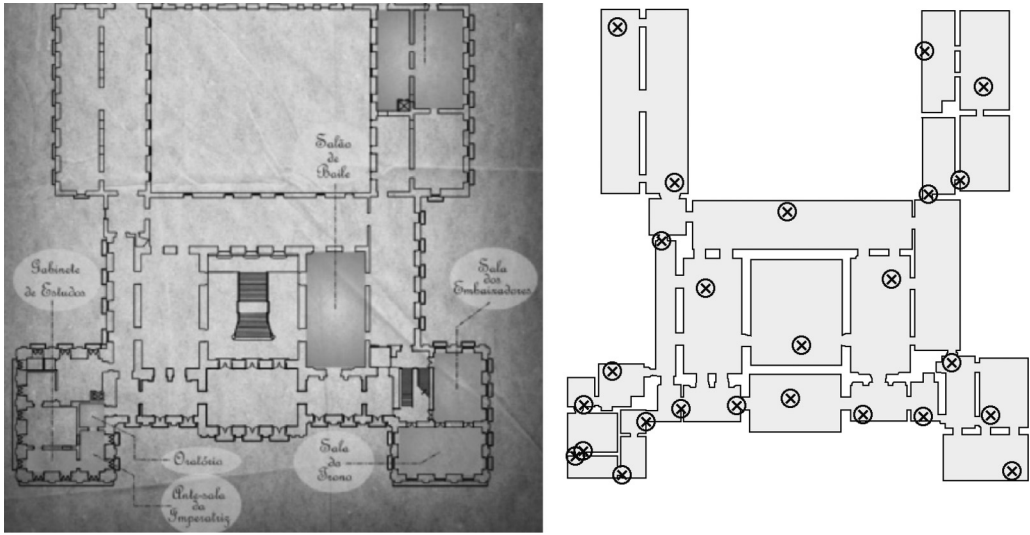


Fig. 1. National Museum floor plan (left); an optimal guard positioning (right).

it turned into one of the most investigated problems in computational geometry. The breadth of this research can be perceived from the many important works that appeared, including O'Rourke's classical book [O'Rourke 1987], a recent text by Ghosh on visibility algorithms [Ghosh 2007], and surveys by Shermer [1992] and Urrutia [2000].

In order to illustrate the problem, consider the Brazilian National Museum of Natural History and Anthropology, in Rio de Janeiro. Figure 1 (left)¹ shows the floor plan of its second floor. Suppose the museum's curator intends to place cameras to guard this entire floor while subject to a limited budget. What would be the smallest number of cameras and their positions to completely cover this area? Figure 1 (right) shows an optimal solution that uses 25 cameras.

Besides the natural application just described, the study and resolution of the AGP and its variations may also be of value in other areas. Consider, for instance, an application where one seeks to place the smallest number of nodes forming a sensor network that covers an area under the restriction that the viewing range is limited (in angle or in depth). Similarly, a robot equipped with a 3D scanner (of bounded range) may be used to map out a building floor, a plaza, or an entire town. Minimizing the number of scan positions is also a straightforward variation of the AGP (see Borrmann et al. [2013]).

From these applications, it is easy to see that by varying the concept of visibility, changing what constitutes coverage, or predefining a discrete set of guard candidates, we are led to a number of variations of the AGP. In the *Art Gallery Problem With Fixed Guard Candidates* (AGPFC), in particular, a viable solution consists of a set of guards that guarantee surveillance of the entire polygon while having their placement restricted to a finite number of positions in the polygon. In contrast, consider the *Art Gallery Problem With Witnesses* (AGPW), which asks for a set of guards able to observe merely a given finite set of points inside the polygon, while guard placement may be unrestricted. Both of these versions as well as many others are known to be NP-hard [O'Rourke 1987].

¹Obtained from www.museunacional.ufrj.br/visitacao/redescobrimdo-a-casa-do-imperador/pavimento-2.

In the past 10 years, advances have been reached in the development of algorithms for solving the AGP in practice. Among them, we find approximation algorithms [Ghosh 2010], heuristics [Amit et al. 2007; Bottino and Laurentini 2011], and, more recently, some advances in the search for exact algorithms [Couto et al. 2011; Kröller et al. 2012; Fekete et al. 2015]. However, none of the algorithms proposed for the original AGP (where guards may be placed freely inside the polygon) was able to consistently find optimal solutions for large instances within a reasonable amount of time.

In this article, we describe a practical algorithm for solving the original AGP to exactness. A first version of the algorithm was outlined in Tozoni et al. [2013]. It follows an iterative process of obtaining lower and upper bounds that lead to an optimal solution. These bounds are drawn up through the resolution of instances of the aforementioned AGPW and AGPFC problems using ILP techniques. As we shall see in Section 5, our technique was implemented and tested on standard desktop PCs and optimal solutions were obtained (in a reasonable amount of time) for almost all instances of the benchmark. These results consist of a significant improvement over the previously published techniques. Not only did the proposed algorithm solve instances with up to 2,500 vertices, something without precedent in the literature, but also it proved to be quite robust, achieving optimality in more than 95% of the 2,700+ instances tested. A software package with an implementation of the algorithm for Linux operating systems is provided as supplementary material.

In the next section, we briefly survey relevant related works. Section 3 is dedicated to presenting a few basic concepts and notations, while Section 4.5 describes the main steps of the proposed algorithm. A discussion of the computational results obtained in our experiments appears in Section 5. Lastly, in Section 6, we draw some conclusions and identify future directions of research.

2. RELATED WORK

As early as 1975, Chvátal proved that $\lfloor n/3 \rfloor$ guards are always sufficient and sometimes necessary for covering a simple polygon of n vertices. Shortly thereafter, specialized versions of the problem started being investigated in which the localization of the guards is restricted to prespecified positions. Thus, the concepts of vertex and edge guards arose to refer to the cases where guards can only be placed at vertices or set on edges of the polygon, respectively. The term “point guards” came to be used to characterize the situation where guard positioning can freely take place anywhere in the interior (or on the boundary) of the polygon. About a decade after Chvátal’s result, Lee and Lin proved the NP-hardness of the AGP for vertex guards [Lee and Lin 1986]. The cases of point guards and of edge guards are also known to be equally hard [O’Rourke 1987].

After progress on theoretical results regarding the complexity of the AGP, researchers turned their attention to the development of algorithms for positioning guards and trying to reach solutions as close as possible to optimality. Ghosh, in 1987, proposed a $\log n$ -approximation algorithm for the AGP with vertex or edge guards [Ghosh 1987]. The idea, also exploited more recently for exact algorithms, was to reduce the AGP to the *Set Cover Problem* (SCP) and to apply an already known approximation for the SCP.

In 2007, Amit et al. [2007] proposed a set of heuristics based on greedy strategies and polygon partition methods. Four years later, Bottino and Laurentini [2011] proposed a new heuristic in which an initial solution for the AGP is computed from an optimal solution to the related problem of covering only the edges of the polygon. Both algorithms were a step forward and the latter, in particular, was able to find optimal solutions for a considerable number of instances.

Recently, however, the search for algorithms that are able to find provably optimal solutions for specific formulations has intensified. In this line, two results are worth

highlighting. In 2011, Couto et al. [2011], published the first exact algorithm for the Art Gallery Problem with Vertex-Guards (AGPVG) with verified practical efficiency. They devised an iterative technique that was guaranteed to converge to an optimal solution. While the maximum number of iterations was proven to be polynomial in the number of vertices in the worst case, actual convergence happened after just a few iterations. The first interesting insight was to show that the infinite points of the polygon to be watched could be discretized into a finite set of points, called *witnesses*. The AGPVG, having a finite number of possible guards, coupled with the now finite number of witnesses to be covered, was reduced (in the Karp sense) to a sequence of SCP instances that could be solved to optimality using an Integer Programming solver. These instances started off with a small subset of witnesses and simply included additional witnesses whenever the previous solution was not viable for the original instance. An alternative approach consisted of including all predetermined witnesses in a single SCP instance and solving the problem to exactness in a single shot. In fact, in Couto et al. [2011], a clever reduction of the witness set is also presented, which makes this alternative approach much more competitive timewise. These algorithms for the AGPVG were tested on more than 10,000 polygons of various classes and, for up to 2,500 vertices, optimal solutions were attained in just a few minutes of computing time.

Through a related approach, Kröller et al. [2012] tackled the original AGP by seeking to reduce the infinite number of witnesses and guard candidates into finite sets, thus generating instances of the discretized AGP, which can be reduced to SCP instances as mentioned earlier. By solving the linear relaxation of the primal and dual formulations of the SCP instances created, their method is able to find upper and lower bounds for the original problem. The method is also iterative and adds new witnesses and guards at each new iteration, in the hope that an integer solution for the relaxed AGP is found. A negative result is inferred after a pre-established execution time limit is exceeded. Although their algorithm converged in several instances, giving rise to an increase in the number of instances for which optimal solutions are known, unfortunately, for many polygons the method was unable to converge to integer solutions.

This brief summary of the state of the art on exact solutions for the original AGP sets the stage for the contributions that this article brings forth.

3. TERMINOLOGY AND FORMULATION

Recall that the objective of the Art Gallery Problem is to find the minimum number of guards that are sufficient to ensure visual coverage of a given gallery, represented as a simple polygon P (possibly with holes). Consider that the vertex set V of P has cardinality n .

Two points in P are *visible* to each other if the line segment that joins them does not intersect the exterior of P . The *visibility polygon* of a point $p \in P$, denoted by $\text{Vis}(p)$, is the set of all points in P that are visible from p .

Given a finite set S of points in P , a *covered* (*uncovered*, respectively) region induced by S in P is a maximal connected region in $\cup_{p \in S} \text{Vis}(p)$ ($P - \cup_{p \in S} \text{Vis}(p)$, respectively).

Moreover, the union of the edges in $\text{Vis}(p)$ over all points p in S defines a geometric arrangement, denoted $\text{Arr}(S)$, that partitions P into a collection of closed polygonal faces called *Atomic Visibility Polygons* or simply *AVPs*. Denote by $C_f \subset S$ the set of points in S that cover an AVP f . Define a partial order $<$ on the faces of $\text{Arr}(S)$ as follows. Given two AVPs f and f' , we say that $f < f'$ if they are adjacent and $C_{f'} \subset C_f$. We call f a *shadow face* (*light face*) if f is locally minimal (maximal) with respect to $<$. Figures 2 and 3 illustrate these concepts. Local minimality and maximality can be understood from Figure 3 (right). It is easy to see by inspection that the white regions are covered by exactly one guard, whereas the light gray ones are covered by two guards. Besides, no guard is able to see a dark gray region. Therefore, according to the




Fig. 2. A polygon with holes (left), the visibility polygon of a point (center); the arrangement induced by a finite set S of points (right).




Fig. 3. A set S of three points and its covered (gray) and uncovered (white) regions (left); the arrangement induced by S with the *light* (gray) and *shadow* (dark gray) AVPs (right).

previous definitions, the dark (light) gray regions are local minima (maxima) relative to \prec .

One important observation for the complexity analysis of our algorithm is that, for a given set $S \subseteq P$, the time needed to build $\text{Arr}(S)$, as well as the number of AVPs in this arrangement, is $O((|S| + |P|)^3)$ [Bose et al. 2002], where $|P|$ denotes the number of vertices in P (i.e., $|P| = n$).

3.1. ILP Formulation

In a geometric setting, the AGP can be restated as the problem of determining a smallest set of points $G \subset P$ such that $\cup_{g \in G} \text{Vis}(g)$ equals P . This leads to a reduction from the AGP to the SCP, in which the points in P are the elements to be covered and the visibility polygons of the points in P are the sets used for covering. Accordingly, we can use this reduction to construct an ILP formulation for the AGP:

$$\min \sum_{c \in P} x_c \quad (1)$$

$$\text{s.t.} \quad \sum_{\substack{c \in P \\ w \in \text{Vis}(c)}} x_c \geq 1, \forall w \in P \quad (2)$$

$$x_c \in \{0, 1\}, \forall c \in P. \quad (3)$$

In this model, a binary variable x_c is assigned to each point c in P and is set to 1 only when a guard is positioned at c in the optimal solution. Accordingly, the objective




Fig. 4. An illustration of the four different variants of the Art Gallery Problem: (a) AGP, (b) AGPFC, (c) AGPW, and (d) AGPWFC. The witnesses and the guard candidates are identified by the symbols “ \times ” and “ \otimes ”, respectively. Darker guard candidates refer to guards present in an optimal solution of the corresponding problem and, when appropriate, have their visibility polygons also depicted.

function (Equation (1)) seeks to minimize the number of guards, whereas the inequalities (Equation (2)), the *coverage constraints*, impose that each point in P must be seen by at least one guard. Although correct, this formulation has an infinite number of constraints and an infinite number of variables, rendering it of no practical benefit. A natural idea that arises is to make at least one of these quantities finite. By fixing only the guard candidates to be a finite set, we obtain the so-called *Art Gallery Problem With Fixed Guard Candidates*. On the other hand, by restricting solely the witness set to a finite set, we end up with the special AGP variant known as the *Art Gallery Problem With Witnesses*. In principle, in the first case, we are still left with an infinite number of constraints, while in the second case, we still have an infinite number of variables. However, in order to have a tractable SCP instance, one should have both the witness and the guard candidate sets of finite size. The AGP variant that fulfills this property is named the *Art Gallery Problem with Witnesses and Fixed Guard Candidates* (AGPWFC). Examples of these three versions of the AGP are shown in Figure 4.

A remarkable theoretical result discussed in the next section is that any instance of the AGPW or of the AGPFC can be reduced in polynomial time to an instance of the AGPWFC. This means that, to solve the AGPW or the AGPFC, one actually has to find an optimal solution of a polynomial-sized SCP instance. Now, notice that the optimum of an AGPW instance always yields a lower bound for the original AGP value, since it is clearly a relaxation of the latter. On the other hand, if the set of guard candidates covers the polygon P , an optimal value for the AGPFC is feasible for the original AGP instance and, therefore, provides an upper bound for it. These are relevant observations since modern ILP solvers are capable of handling very large-sized SCP instances such as AGPWFC instances, for that matter. As we will see later, our algorithm relies on these facts.

To close this section, a couple of additional notations are introduced. Let D and C be finite witness and guard candidate sets, respectively. We denote the AGPW, AGPFC, and AGPWFC problems defined for the sets C and D by $\text{AGPW}(D)$, $\text{AGPFC}(C)$, and $\text{AGPWFC}(D, C)$, respectively. The SCP model associated to $\text{AGPWFC}(D, C)$ is then

$$\min \sum_{c \in C} x_c, \quad (4)$$

$$\text{s.t.} \quad \sum_{\substack{c \in C \\ w \in \text{Vis}(c)}} x_c \geq 1, \forall w \in D, \quad (5)$$

$$x_c \in \{0, 1\}, \forall c \in C. \quad (6)$$

In this model, a binary variable x_c is assigned to each point c in C and is set to 1 only when a guard is positioned at c in the optimal solution. As in the infinite case, the objective function (Equation (4)) minimizes the number of guards, whereas inequalities (Equation (5)) enforce that each witness point in D be seen by at least one guard. Since C and D are finite, this formulation has a finite size.

4. A PRACTICAL ITERATIVE ALGORITHM FOR AGP

The purpose of this section is to describe our algorithm, which basically consists of iteratively obtaining increasing lower (dual) and decreasing upper (primal) bounds for the AGP. These steps are repeated until either the gap between the bounds reaches zero or the time limit is exceeded. New bounds are obtained by solving discretized versions of the original AGP. To find a new lower bound, an AGPW instance is solved, while in the upper-bound case, an AGPFC instance is worked out in which the guard candidate set covers the polygon. One important fact to remember is that an optimal solution for such an instance of AGPFC is also viable for AGP, since the AGPFC asks for a solution that guards the entire polygon. Consequently, reducing the gap to zero means reaching an optimal solution for the original AGP.

In Algorithm 1, we present a simplification of how our technique works. After initializing the witness and guard candidate sets, the algorithm enters its main loop (lines 4–10). At each iteration, new lower and upper bounds are computed and, if optimality has not been proved, the witness and guard candidate sets are updated in line 8. Later we will see how this can be done in a way in which the duality gap decreases monotonically through the iterations. It is also worth noticing that, as we still do not have a proof of convergence for the algorithm, the parameter MAXTIME is used to limit its running time.

ALGORITHM 1: AGP Algorithm

```

1: Set  $UB \leftarrow |V|$  and  $LB \leftarrow 1$ 
2: Select the initial witness set  $D$ 
3: Select the initial guard candidate set  $C$ , so that  $V \subset C$ 
4: while ( $UB \neq LB$ ) and (MAXTIME not reached) do
5:   Solve  $\text{AGPW}(D)$ , set  $G_w \leftarrow$  optimal solution of  $\text{AGPW}(D)$  and  $LB \leftarrow \max\{LB, |G_w|\}$ 
6:   Solve  $\text{AGPFC}(C)$ , set  $G_f \leftarrow$  optimal solution of  $\text{AGPFC}(C)$  and  $UB \leftarrow \min\{UB, |G_f|\}$ 
7:   if ( $UB \neq LB$ ) then
8:     Update  $D$  and  $C$ 
9:   end if
10: end while
11: return  $G_f$ 

```

In the following two subsections, the procedures for solving AGPW (line 5) and AGPFC (6) instances are described in detail. Subsequently, Section 4.3 describes the resolution method for AGPWFC instances through ILP techniques. As said in the previous section, both the AGPW and the AGPFC can be cast as an AGPWFC, justifying why we focus on the latter. In Section 4.4, we present how the management of the witness set and, as a consequence, of the guard candidate set is done. Lastly, Section 4.5 gathers all the algorithmic information presented previously and describes the complete algorithm for the AGP.

4.1. Solving AGPW

Consider $\text{AGPW}(D)$ the Art Gallery Problem instance where the objective is to cover all points in a finite set D . As said before, the resolution of an AGPW on D allows for the discovery of a new lower bound for the AGP since fully covering P requires at least as many guards as the minimum sufficient to cover the points in D .

However, despite being a simplification of the original AGP problem, we are still dealing with an infinite number of potential guard positions, which does not allow for a direct reduction to the Set Cover Problem. Thus, our approach is based on discretizing the guard candidate set, creating an AGPWFC instance from our original AGPW. To do this, we apply Theorem 4.1.

THEOREM 4.1. *Let D be a finite subset of points in P . Then, there exists an optimal solution for $\text{AGPW}(D)$ in which every guard belongs to a light AVP of $\text{Arr}(D)$.*

PROOF. Let G be an optimal (cardinality-wise) set of guards that covers all points in D . Suppose there is a guard g in G whose containing face F in $\text{Arr}(D)$ is not a light AVP. This means that F is not locally maximal with respect to $<$ (see Section 3). In other words, there exists a face F' of $\text{Arr}(D)$ that shares an edge with F such that $F < F'$; that is, F' sees more points of D than F does. An inductive argument suffices to show that this process eventually reaches a light AVP (locally maximal w.r.t. $<$) wherein lies a point that sees at least as much of D as g does; that is, g may be replaced by a guard that lies on a light AVP. The theorem then follows, by induction, on the number of guards of G . \square

From Theorem 4.1, one concludes that there exists an optimal solution for $\text{AGPW}(D)$ in which all the guards are in light AVPs of the arrangement induced by D . Besides, as every point (whether inside or on the boundary) of an AVP can see the same set of witnesses, we can state that there is an optimal solution G where each point in G belongs to the set $V_{\mathcal{L}}(D)$ of all vertices from the light AVPs of $\text{Arr}(D)$. Therefore, an optimal solution for $\text{AGPW}(D)$ can be obtained simply by solving $\text{AGPWFC}(D, V_{\mathcal{L}}(D))$. As seen before, the latter problem can be modeled as an ILP, where the numbers of constraints and of variables are polynomial in $|D|$. This follows, as mentioned in Section 3, from the fact that the number of AVPs (and vertices) in $\text{Arr}(D)$ is bounded by a polynomial in $|D|$ and $|P|$. Algorithm 2 shows a pseudo-code of the AGPW resolution method.

ALGORITHM 2: $\text{AGPW}(D)$ Algorithm

- 1: $\text{Arr}(D) \leftarrow$ construct the arrangement from the visibility polygons of the points in D
 - 2: $V_{\mathcal{L}}(D) \leftarrow$ identify the vertices of the light AVPs of $\text{Arr}(D)$
 - 3: $C \leftarrow V_{\mathcal{L}}(D)$
 - 4: Solve $\text{AGPWFC}(D, C)$: set $G_w \leftarrow$ optimal solution of $\text{AGPWFC}(D, C)$
 - 5: **return** G_w
-

4.2. Solving AGPFC

As we now know how to generate dual bounds for the AGP, the next task is to compute good upper bounds for the problem. A possible way to achieve this is through the resolution of an AGPFC instance in which the guard candidate set is known to cover the polygon. As said earlier, under this condition, an AGPFC solution is always viable for the original problem.

In this context, consider $\text{AGPFC}(C)$ the instance of AGPFC where the finite guard candidate set is C . In contrast to what was discussed regarding the AGPW, we now have a finite number of guard candidates and an infinite number of points to be covered. Therefore, a straightforward resolution method using a reduction to SCP is not possible. To circumvent this, our algorithm discretizes the original AGPFC instance, relying on what Theorem 4.2 establishes.

THEOREM 4.2. *Let D and C be two finite subsets of P , so that C fully covers P . Assume that $G(D, C)$ is an optimal solution for $\text{AGPWFC}(D, C)$. If $G(D, C)$ also covers P , then $G(D, C)$ is an optimal solution for $\text{AGPFC}(C)$.*

PROOF. Assume that $G(D, C)$ covers P , but it is not an optimal solution for $\text{AGPFC}(C)$. Then, there exists $G' \subseteq C$ with $|G'| < |G(D, C)|$ such that G' is a feasible solution for $\text{AGPFC}(C)$; that is, G' covers P . This implies that G' is also a feasible solution for $\text{AGPWFC}(D, C)$, contradicting the fact that $G(D, C)$ is optimal for this problem. \square

Theorem 4.2 shows that an optimal solution for AGPFC may be obtained through the resolution of an AGPWFC instance, provided it fully covers P . Whenever an optimal solution for the simplified version (AGPWFC) leaves uncovered regions in P , additional work is required. To guarantee that we attain an optimal solution for AGPFC, we will employ here a technique designed by Couto et al. [2011] to solve the AGPVG, a special case of AGPFC where $C = V$, but which may be used to solve the general case without significant changes.

Initially, consider that we have a finite witness set D . Using the guard candidate set C , we can create and solve the $\text{AGPWFC}(D, C)$ instance. If the solution fully covers the polygon, we have satisfied the hypothesis of Theorem 4.2 and, consequently, we have an optimal solution for AGPFC. Otherwise, there are regions of the polygon that remain uncovered. We now update the witness set, adding new points within the uncovered regions, denoted $C_u(C)$, and repeat the process.

As shown in Couto et al. [2011], the iterative method for solving AGPFC converges in polynomially many iterations. To clarify this point, consider the arrangement induced by all visibility polygons of guard candidates in C . This arrangement partitions the polygon into $O(|C| \cdot n^2)$ faces (AVPs). These faces have the property that if any of its point is covered, the entire face is covered. Therefore, our aim is to iteratively construct D by choosing witnesses in the interior of AVPs of $\text{Arr}(C)$, leading to an AGPWFC instance whose optimal solution will also be viable and optimal for the AGPFC. While the number of iterations is clearly in $O(|C| \cdot n^2)$, in Couto et al. [2011] it is shown that, in practice, this number is much lower. Moreover, it can even be argued that it suffices to select one point from each shadow AVP of the arrangement induced by C . See Couto et al. [2011] for details.

A pseudo-code for the algorithm used to solve the AGPFC is shown in Algorithm 3.

4.3. Solving AGPWFC (SCP)

Having reduced both AGPW and AGPFC problems into AGPWFC instances in order to obtain the desired bounds, the objective becomes solving the latter efficiently. Since an $\text{AGPWFC}(D, C)$ instance can easily be reduced to an SCP, where the witnesses in D are the elements to be covered and the visibility polygons of the guard candidates in C

ALGORITHM 3: AGPFC(C) Algorithm

```

1:  $D_f \leftarrow$  initial witness set
2: repeat
3:   Solve AGPWFC( $D_f, C$ ): set  $G_f \leftarrow$  optimal solution
4:   if  $G_f$  does not fully cover  $P$  then
5:      $D_f \leftarrow D_f \cup C_{\mathcal{U}}(G_f)$ 
6:   end if
7: until  $G_f$  fully covers  $P$ 
8: return  $G_f$ 

```

are the subsets considered, we will make use of the ILP formulation for SCP presented in Section 3.

A simple and straightforward approach would be to directly use an ILP solver, such as XPRESS or CPLEX, since even large instances of the (NP-hard) SCP can be solved quite efficiently by many modern integer programming solvers. However, as observed in our experiments, some AGP instances can generate significantly complex and very large SCP instances, rendering the solvers less efficient. For this reason, some known techniques were implemented to improve the solvers' running time. Among these, the most effective consisted of the reduction on the number of constraints and variables in the initial model. A simple primal heuristic based on Lagrangian relaxation was also implemented to obtain viable solutions. Despite the fact that these are standard techniques to ILP practitioners, some of them proved to be instrumental in the context of the AGP. So, for completeness, we briefly describe them next.

The method used for reducing constraints and variables is based on containment relationships between columns (guard candidates) and between rows (witnesses) of the Boolean constraint matrix corresponding to the ILP model of the SCP instance.

First, we search for and discard *redundant guard candidates* (columns). A guard candidate g_r is *redundant* if there is another candidate that covers the same witness set as g_r . For this step, we divide all guard candidates into groups, according to the light AVP they belong to. Guard candidates from the same AVP are then tested pairwise for redundancy, and, when one is found, it is removed from the original matrix. Here a slight improvement applies to SCP instances arising from the AGPW. Remember that when lower bounds are computed, the witness set D remains fixed while the guard candidates are the vertices of the light AVPs in $\text{Arr}(D)$. In this case, by definition, all vertices of a given light AVP cover the same subset of witnesses. Hence, all but one of the corresponding variables is redundant in the SCP instance; that is, we are left with a single vertex from each light AVP. This reduces the size of the guard candidate set by a factor of at least three, although the remaining subset still needs to be checked pairwise for redundancy.

After the removal of columns, we also test each pair of rows in search for removable ones. We say that a row represents an *easy witness* w whenever the set of guard candidates that see w properly contains the (whole) set of candidates that see some other witness.

It is important to notice that the test for redundant candidates (columns) and easy witnesses (rows) can be performed simply by using very fast bit-string operations.

Besides conducting matrix reduction, our algorithm also uses an initial *Lagrangian Heuristic* (LH) with the goal of finding good (or optimal) viable solutions for the AGP-WFC. The heuristic implemented is based on the work by Beasley [1993], which, among other results, presents a *Lagrangian Relaxation* for SCP. The idea of the relaxation is to move all coverage constraints into the objective function and use penalties u_w (for each constraint), the Lagrangian Multipliers (LMs), resulting in the following *Lagrangian*

Primal Problem (LPP):

$$z(u) = \min \sum_{c \in C} x_c + \sum_{w \in D} u_w \left(1 - \sum_{\substack{c \in C \\ w \in \text{Vis}(c)}} x_c \right)$$

$$x_c \in \{0, 1\}, \forall c \in C.$$

After this step, the LPP (with no actual constraints) can be solved by inspection. It is well known that the optimum of this problem gives a lower bound for the original SCP instance. The new quest is to find the values for LMs (u_w variables) that maximize this lower bound. The optimization problem of finding the best Lagrangian Multipliers is called the *Lagrangian Dual Problem (LDP)*:

$$Z = \max z(u)$$

$$u_w \geq 0, \forall w \in W.$$

A classical approach to tackle the LDP is to apply the *subgradient method* (c.f., Beasley [1993]), which is an algorithm in which the LM values are updated iteratively using a subgradient of the objective function. In this way, at each iteration, a new LPP instance is solved and, based on this result, we apply a primal heuristic, which attempts to find a new viable solution for the original SCP instance. The primal heuristic consists of a greedy algorithm that uses the Lagrangian costs obtained during the last LPP resolution to define which candidates will be selected to join the new viable solution (see Beasley [1993] for more details). In summary, at each iteration, a new lower bound for the SCP is obtained from the LPP resolution and a new upper bound is obtained by the primal heuristic. The subgradient method stops when either an optimal solution is found or a maximum number of iterations has been reached.

After running the subgradient method for solving LDP, we have a reduced ILP matrix and an upper bound for AGPWFC. This information can be used in the ILP Solver to improve its performance. Algorithm 4 summarizes the steps for solving the AGPWFC instance, using matrix reduction and a Lagrangian heuristic.

ALGORITHM 4: AGPWFC(D, C) Algorithm

- 1: $M \leftarrow$ Boolean matrix of the SCP model constructed using D and C
 - 2: $R \leftarrow$ matrix obtained from M after applying the reduction procedure
 - 3: Set $I \leftarrow$ SCP instance associated to R
 - 4: Solve I using LH
 - 5: Set $\text{LB}_{\text{SCP}} \leftarrow$ best lower bound found for I by the subgradient method
 - 6: Set $G_v \leftarrow$ best viable solution found for I by the primal heuristic
 - 7: $G_s \leftarrow G_v$
 - 8: **if** $|G_s| \neq \text{LB}_{\text{SCP}}$ **then** {LH was not able to prove optimality}
 - 9: Solve I using an ILP Solver with the primal bound G_v
 - 10: Set $G_s \leftarrow$ optimal solution found for the ILP model
 - 11: **end if**
 - 12: **return** G_s
-

4.4. Witness Management

The witnesses selected during the execution of our algorithm play an important role in the search for an optimal solution for the AGP. This occurs because the choice of the witness set D directly affects the quality of the lower and upper bounds and, consequently, the resolution of the AGPW (as well as the AGPFC) instances, as we will see later.

Recall that, before the first iteration of Algorithm 1, an initial witness set is chosen and in the following ones, it gets suitably updated (line 5). Several initialization alternatives were considered, tested, and analyzed. We present next the three most effective of them.

The first initialization choice, called *All-Vertices* (AV), consists of using all vertices of the polygon as witnesses, that is, $D = V$. Second, in an attempt to start off with a small number of witnesses, we also considered initializing D with only the convex vertices of the polygon, which is referred to as the *Convex-Vertices* (CV) initialization. The reason for trying to reduce the initial witness set lies in the fact that a smaller such set is likely to lead to a lower number of visibility polygon calculations. Moreover, the visibility arrangement used becomes less complex and so does the SCP model.

The third and last alternative is based on a work by Chwa et al. [2006], where the authors define a polygon P to be *witnessable* when there exists a finite witness set $W \subset P$ with the property that any set of guards that covers W must also cover the entire polygon. That paper also presents an algorithm that computes a minimum witness set for a polygon whenever it is witnessable. The method for constructing this minimum witness set consists of placing a witness in the interior of every *reflex-reflex* edge of P and on the convex vertices of every *convex-reflex* edge. The terms *convex* and *reflex* here refer to the interior angles at a vertex or at the endpoints of an edge. Based on these selection criteria, we devised our third discretization method, called *Chwa-Points* (CP), which assembles the initial witness set for our algorithm from the midpoints of all reflex-reflex edges and all convex vertices from convex-reflex edges.

It follows from the results in Chwa et al. [2006] that, when the Chwa-Points discretization is used for a witnessable input polygon, our AGP algorithm will find an optimal solution in a single iteration. However, as observed in our experiments, non-witnessable polygons are the norm. In fact, among our random benchmark instances, they constitute the vast majority.

Let us now focus on the updating process of the witness set throughout the algorithm iterations. Of course, the better the choice of a new set of points for inclusion into the witness set, the better the lower or the upper bound obtained would be. In essence, the process consists simply of adding points from the uncovered regions arising from the solution of the previous AGPW instance.

Our experiments showed that the inclusion of an interior point from each uncovered region was not sufficient to lead the algorithm to good performance and convergence. The shortfall was traced to the absence of new points on the boundary of the polygon, which proved to be useful to the evolution of the bounds obtained. Therefore, whenever an edge of an uncovered region is found to be contained on the boundary of the polygon, its vertices and its midpoint were also selected to increment the witness set throughout the iterations. These points along with an interior point of each uncovered region formed the whole set M of new witnesses.

4.5. Resulting Algorithm

Once each of the main steps of the algorithm is understood, we are able to describe how these parts fit together to compose the complete algorithm. Algorithm 5 sums up how the process as a whole works.

It is important to notice that the set of guard candidates used in the AGPW resolution, which consists of V plus all vertices from the light AVPs induced by D , is actually the set C from the AGPFC(C) instance solved on line 14. This means that the AGPW resolution is the first iteration of the AGPFC algorithm (Algorithm 3). Thus, all information obtained during the solution of AGPW(D) can be reused for AGPFC(C), which improves the performance of the implementation.

ALGORITHM 5: AGP Algorithm

```

1:  $D \leftarrow$  initial witness set {See Section 4.4}
2: Set  $LB \leftarrow 1$ ,  $UB \leftarrow n$  and  $G^* \leftarrow V$  {initialize lower (LB) and upper (UB) bounds &
   best-known solution}
3: loop
4:   Solve AGPW( $D$ ): set  $G_w \leftarrow$  optimal solution and  $z_w \leftarrow |G_w|$  {Section 4.1: LB computation
   with set  $D$ }
5:   if  $G_w$  is a coverage of  $P$  then
6:     return  $G_w$  { $G_w$  is feasible:  $z_w$  is both an LB and a UB}
7:   end if
8:    $LB \leftarrow \max\{LB, z_w\}$  {update best-known LB}
9:   if  $LB = UB$  then
10:    return  $G^*$  {optimal solution found}
11:   end if
12:    $C \leftarrow V_{\mathcal{L}}(D) \cup V$  {guard candidates for UB computation: vertices of  $P$  and light AVPs of
    $\text{Arr}(D)$ }
13:    $U \leftarrow C_{\mathcal{U}}(G_w)$  {new witnesses for UB computation: one additional point per uncovered
   region}
14:    $D_f \leftarrow D \cup U$  {enlarged witness set for UB computation}
15:   Solve AGPFC( $C$ ), using  $D_f$ : set  $G_f \leftarrow$  optimal solution and  $z_f \leftarrow |G_f|$  {Section 4.2: UB
   computation}
16:    $UB \leftarrow \min\{UB, z_f\}$  and, if  $UB = z_f$  then set  $G^* \leftarrow G_f$  {update best known UB and
   solution if needed}
17:   if  $LB = UB$  then
18:     return  $G_f$  {optimal solution found}
19:   end if
20:    $D \leftarrow D \cup U \cup M$  {update witness set for next iteration; see Section 4.4 for the definition
   of  $M$ }
21: end loop

```

Another relevant aspect of this algorithm is that information on bounds may be used throughout the iterations in order to skip unnecessary steps. For instance, if an upper bound UB was found in a previous iteration and a new AGPFC instance is being solved, whose current solution is not lower than UB , then we may stop the AGPFC resolution before obtaining an optimal solution since the upper bound cannot be improved.

Figures 5, 6, and 7 illustrate the execution of the AGP algorithm on an orthogonal polygon.

5. COMPUTATIONAL RESULTS

To report how the experiments, conducted to validate our algorithm, were performed, we describe the computing environment, the instances used, the parameters employed, and how they affect the overall performance of the resulting implementation. Lastly, a comparison with other existing techniques is presented, followed by a summary of the results obtained with all tested instances.

5.1. Computing Environment

All the experiments were performed on standard PCs featuring an Intel Core i7-2600 at 3.40GHz and 8GB of RAM, and running Ubuntu 12.04.1 LTS (GNU/Linux 3.2.0-31-generic x86_64).

Our algorithm (Algorithm 5) was coded in the C++ programming language. The program uses the Computational Geometry Algorithms Library (CGAL) [CGAL 2012], version 3.9, to benefit from visibility operations, arrangement constructions, and other




Fig. 5. **Solving the AGPW (lower bound):** The initial witness set D (Chwa-Points) (left); the arrangement $\text{Arr}(D)$ and the light AVPs (center); the solution to $\text{AGPW}(D)$ (right).




Fig. 6. **Solving AGPFC (upper bound):** Updated witness set D_f (left); the solution to $\text{AGPWFC}(D_f, V_L(D) \cup V)$ and to $\text{AGPFC}(V_L(D) \cup V)$ (right).




Fig. 7. **Solving the AGPW (lower bound):** The new witness set D (left); the arrangement $\text{Arr}(D)$ and the light AVPs (center); the solution to $\text{AGPW}(D)$ and to the AGP (right).

geometric tasks. To solve the integer programs that model the SCP instances, we used the XPress Optimization Suite [XPress 2013], version 7.0.

To achieve accurate time measurements, all tests were run in isolation; that is, no other process was being executed concomitantly. Moreover, the experiments were run on a single core of the processor, hence not profiting from parallelism. Each process was given a time limit of 60 minutes, after which the instance was considered unsolved and the program was terminated.

5.2. Instances

In total, our benchmark included more than 2,700 publicly available test polygons with and without holes. By collecting such an extended experimentation testbed, from various sources, composed of polygons of multiple classes and sizes, we were able to stress the algorithm robustness.

For polygons without holes, three classes of polygons were used: simple, orthogonal, and Von Koch polygons. The sets of simple polygons were obtained from the work of Bottino and Laurentini [2011] (250 polygons) and Couto et al. [2011] (600 polygons).

Random orthogonal polygons included 80 instances from Bottino and Laurentini [2011] and another 600 instances from Couto et al. [2011], while a third set of hole-free polygons contained 480 Von Koch polygons also from Couto et al. [2011].

The instances were grouped according to their sources and sizes: those from Bottino et al. contain between 30 and 60 vertices and are subdivided into groups of 20 polygons each. Those from Couto et al. range from 20 to 2,500 vertices and each group has 30 instances of equal size.

In the case of polygons with holes, 120 instances called *spike polygons* ranging with 60, 100, 200, and 500 vertices were obtained from the work by Kröller et al. [2012]. As we point out later, the availability of these instances afforded us the possibility of comparing our method with the technique presented in Kröller et al. [2012], which currently is among the most promising methods for solving the AGP.

On the other hand, to make up for the lack of diversity of instances with holes in the available benchmarks, we generated another 600 random polygons with holes. These additional instances were made available on our web page [Couto et al. 2013] so that they can be used for further comparisons.

Given the nature of the polygons created for our benchmark, a description of the generation process is due. We dedicate the next few paragraphs to this purpose. To construct these new polygons with holes, two different generation techniques were used, each responsible for producing half of these new instances.

To understand the process of generating random simple polygons, picture a randomly generated set of points uniformly distributed inside a given rectangle. CGAL's `random_polygon_2` procedure generates a (not necessarily simple) polygon whose vertices are given by a random permutation of those points and applies a method of elimination of edge intersections based on *2-opt moves* (see van Leeuwen and Schoone [1981]).

Now, for the first technique, denoted here *GB*, the process works as follows: after creating the outer boundary for a random simple polygon with holes, consider that we are left with v vertices to be distributed among h holes. After generating a random uniform partition of v into h parts, we iteratively generate the h holes in the following way. At each iteration, we randomly select a point in the interior of the polygon around which we center an isothetic square entirely contained inside the polygon. This square is then stretched in each of the four orthogonal directions, chosen in random order, by λD , where λ is a stretch factor randomly picked from the interval $[0.25, 0.75]$ and D is the maximum elongation within the polygon in that direction. A hole, as a simple




Fig. 8. Instance representing a square in Bremen, Germany (left); an optimal guard positioning (right).

polygon, is then created within this placeholder rectangle, having its number of vertices chosen from one of the unused parts of the previously mentioned random partition of v . Here, for an instance with a total of n vertices, $n/4$ of them were assigned to the outer boundary and $3n/4$ of them distributed among $n/10$ holes.

In contrast, in the second technique, denoted *GD*, the outer boundary of the polygon as well as all the holes are random orthogonal polygons. The maximal placeholder rectangle is constructed so that the chosen random interior point c is one of its vertices. Then, at each iteration, besides randomly selecting a new hole size, two holes are allowed to intersect, in which case they merge into a single hole. Similarly to the previous generator, for an instance with a total of n vertices, the outer boundary has roughly $n/4$ of them, while the remainder are distributed among the resulting $n/10$ holes.

The range $[100, 1000]$, with step size 100, was used for the number of vertices of the polygons with holes. A total of 30 polygons of each size per generator were produced using these methods, resulting in *GB* creating 300 instances (the *simple-simple* class) and an equal number produced by *GD* (the *ortho-ortho* class).

All instances experimented in the work of Couto et al. and Tozoni et al. can be found in our project page [Couto et al. 2013].

The undertaking that led to the improvements in our algorithm is relevant in its own right. It is our belief that it illustrates a process that may guide similar endeavors, so we devote Section 5.3 to its description.

5.3. Development History

The undertaking that led to the improvements in our algorithm is relevant in its own right. It is our belief that it illustrates a process that may guide similar endeavors, so we devote this section to its description.

Since the present work expands and improves our own previous work [Tozoni et al. 2013], a comparison of results is due. Let us take, as an example, the polygon in Figure 8, which is a real instance representing a square in Bremen, Germany, with 14 holes and 311 vertices taken from Borrmann et al. [2013]. More than 20,000 seconds were required to obtain an optimal solution using version V1 of our implementation reported in Tozoni et al. [2013].

Subsequent improved versions of the algorithm and its implementation, leading to version V2, included the following:

- Upgrading the visibility polygon calculation by avoiding computation related to holes that would not benefit the result
- Cutting short the AGPFC procedure whenever the general upper bound found for the AGP had the same cardinality as the current AGPWFC solution, allowing it to skip to the next iteration of Algorithm 5

As an illustration, V2 took less than 10,000 seconds to obtain an optimal solution for the aforementioned Bremen instance.

Next, by reversing the point of view of visibility testing from the perspective of the guards to that of the witnesses (fewer in number, in our approach), we avoided computing many visibility polygons and reused those already required for setting up the arrangement. Moreover, by caching visibility information based on pairs of points already tested, the performance was further improved. Notice that the ILP matrix construction remained the same and changed little on how the algorithm works. This led to version V3, which solved the Bremen instance 5 times faster than V2.

Recall from Section 4.3 that the ultimate version, V4, of the method described in this work includes the use of the Lagrangian heuristic for solving the SCP model, besides the removal of lines and columns from the original SCP matrix.

Moreover, V4 takes advantage of reusable information from previous iterations. For instance, if a lower bound LB for the AGP instance has been established, Algorithm 2 can be halted when solving a new AGPW(D) instance as soon as a primal solution with cardinality LB is found for the corresponding SCP instance in line 4 (e.g., using the Lagrangian heuristic). This follows from the fact that LB was obtained by solving an AGPW(S) instance, for some $S \subset D$. Another interesting situation happens when we are solving an AGPWFC instance inside the iterative algorithm for AGPFC (Algorithm 3). Since in this procedure witnesses are added and never removed, we can guarantee that the solution of the previous AGPWFC instance is a lower bound for the next instance. Using V4, it took less than 600 seconds to reach an optimal solution for the Bremen instance!

5.4. Discretization Techniques for the Witness Set

In Section 4.4, we presented three possible initial discretizations for the witness set. Natural questions that arise are: Which of these leads to the highest number of exact solutions among multiple classes of instances of different sizes? With fewer iterations? With best time performance? Table I and Figure 9 display results obtained in this direction, which exemplify the behavior of the vast majority of the cases observed. In this summary, only ortho-ortho and simple-simple instances with sizes ranging from 100 to 500 vertices were displayed.

In Table I, we have the number of instances that were solved to optimality using each discretization. As said before, the tests were halted after 1 hour of execution time. We can see that *Chwa-Points* and *Convex-Vertices* achieved the best results, solving 294 and 292 of the 300 instances considered, respectively.

By analyzing the charts in Figure 9, which presents the average number of iterations and the average time of successful runs, we see a reasonable advantage when using *Chwa-Points*. In the case of simple-simple polygons with 500 vertices, for example, we have a savings of more than 200 seconds when using *Chwa-Points* as opposed to *Convex-Vertices*. If we compare to the *All-Vertices* technique, this difference grows to more than 500 seconds, on average.

Table I. Number of Instances Solved to Optimality When Using Each of the Three Initial Discretization Techniques

Instance Group	n	# of Instances Solved		
		All-Vertices	Convex-Vertices	Chwa-Points
Simple-simple (30 inst. per size)	100	30	30	30
	200	30	30	30
	300	29	30	30
	400	29	30	28
	500	26	27	28
Ortho-ortho (30 inst. per size)	100	30	30	30
	200	29	29	30
	300	29	29	30
	400	30	30	30
	500	27	27	28



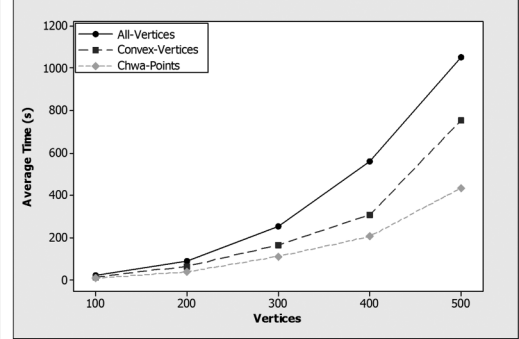

Simple-simple:**Ortho-ortho:**

Fig. 9. Average number of iterations per size (left); average time spent per size (right).

The graphs for All-Vertices and Convex-Vertices show that while they lead to very similar results in regard to the number of iterations, the latter has a considerable advantage over the former in runtime and consequently in the number of instances solved to optimality (Table I). This behavior stems from the fact that the discretization points in the Convex-Vertices strategy is a subset of the ones in the All-Vertices, and clearly a larger witness set leads to more visibility polygon computations, more complex visibility arrangements, and, consequently, larger SCP instances to be solved.

The Chwa-Points discretization technique, having been the most successful one, was chosen as the standard initial discretization for our method. In this context,

Table II. Number of Instances Solved to Optimality When Using or Not Using the Lagrangian Heuristic Method

Instance Groups	n	Instances Solved	
		Without LH	With LH
Simple-Simple (30 inst. per size)	100	30	30
	200	30	30
	300	30	30
	400	28	28
	500	29	28
Ortho-ortho (30 inst. per size)	100	30	30
	200	30	30
	300	29	30
	400	30	30
	500	29	28

throughout the following sections, all results presented were obtained using this initial discretization.

5.5. Lagrangian Heuristic

In Section 4.3, the Lagrangian heuristic was presented as a means to reach optimal solutions for AGPWFC (SCP) instances faster. Lagrangian relaxation is reported (c.f., Beasley [1993]) to perform well in many SCP applications, and therefore, we decided to test it on the SCP instances arising in the context of the AGP. In this section, we briefly analyze the pros and cons of using this technique in our implementation. For this comparison, four groups of instances, with sizes between 100 and 500, were considered: simple and orthogonal instances from Couto et al. [2011] and the new simple-simple and ortho-ortho classes.

Table II shows results containing the number of instances solved when using (not using) the heuristic. Since optimal solutions were always found for all hole-free instances (with or without the heuristic), the table only shows results for instances from the simple-simple and ortho-ortho classes. From these results, we can see that the number of instances solved within our timeframe of 1 hour did not vary considerably. Actually, in the case of the simple-simple group, the version without the heuristic was able to solve one additional instance.

So, at first glance, it does not seem to be advantageous to employ the Lagrangian heuristic. Nevertheless, the time used to compute viable solutions for SCPs using the heuristic is usually negligible with respect to the time spent in the other parts of our implementation. Besides, some instances benefited from using the heuristic, and, in general, it yielded good primal bounds for SCP, which was its primary goal. Thus, in an attempt to understand why the optimality rate was not positively affected by the usage of the heuristic, we investigated this issue a little further. To this end, we looked at the solutions produced by both the heuristic and the ILP solver XPRESS for some instances. It turned out that, although they had equal sizes, these solutions were often not the same, resulting in different uncovered regions as well as different witness and guard candidate sets. As a consequence, the next sequence of SCPs to be solved changed considerably, deeply affecting the time required to find the optimum of the AGP. This observation shows that more research is needed to distinguish among the optimal SCP solutions those that are more promising for our algorithm, that is, that allow it to solve the AGP quicker. For the time being, we turn our attention to the effects of the Lagrangian heuristic in the total running time.

Regarding the average runtime when using the Lagrangian heuristic, the graphs in Figure 10 show a comparison for the case of each of the four groups tested. These

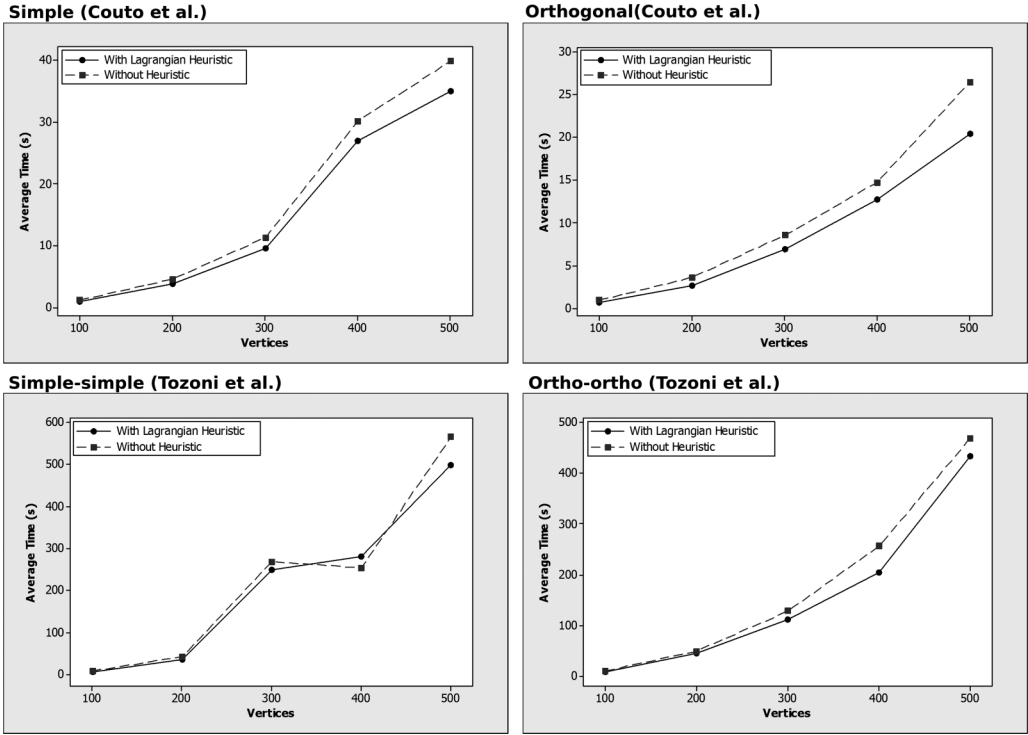


Fig. 10. Comparison between using and not using the Lagrangian heuristic on four different types of polygons.

graphs include only instances that were solved to optimality regardless of whether the Lagrangian heuristic was used.

In general, we perceive a modest advantage when using the Lagrangian heuristic in our procedure. Therefore, we decided to maintain the usage of the Lagrangian heuristic as the default option in all the tests reported in the following sections.

5.6. Comparison with Other State-of-the-Art Techniques

Recently, some techniques achieved good practical results in the search for optimal solutions for the AGP. In 2011, Bottino and Laurentini [2011] presented a heuristic based on an edge-covering algorithm that led to some nice results, while still not able to achieve optimality for a large set of instances. A full comparison between our technique and that heuristic appeared in Tozoni et al. [2013].

Through a rather different approach, Kröller et al. [2012] presented a method that reached optimality for many large and complex instances. For the sake of comparison, we also ran our implementation on the same instances used in that paper. Table III shows the average running times and optimality rates of the two methods. In comparing these results, one has to bear in mind that the computational environments in which the two experiments were conducted were rather different.

It can be seen from Table III that our technique reached optimal solutions for *all* instances within a small average running time. While we initially allowed a time limit of 60 minutes, as opposed to Kröller et al.'s (see Kröller et al. [2012]) 20-minute upper

Table III. Comparison Between the Method of Kröller et al. and Ours

Instance Group	n	Optimality Rates		Average Time (s)	
		Kröller et al. [2012]	Ours	Kröller et al. [2012]	Ours
Simple (30 inst. per size)	60	80%	100%	0.70	0.26
	100	64%	100%	29.40	0.94
	200	44%	100%	14.90	3.77
	500	4%	100%	223.30	35.04
Orthogonal (30 inst. per size)	60	80%	100%	0.40	0.15
	100	54%	100%	1.10	0.64
	200	19%	100%	4.30	2.66
	500	7%	100%	25.30	20.44
Von Koch (30 inst. per size)	60	77%	100%	2.40	0.26
	100	71%	100%	5.50	0.87
	200	70%	100%	18.90	6.63
	500	15%	100%	205.00	76.45
Spike (30 inst. per size)	60	67%	100%	1.70	0.49
	100	68%	100%	3.90	1.40
	200	61%	100%	15.30	7.30
	500	52%	100%	190.20	173.23




Fig. 11. Number of instances solved by our method within discrete amounts of time.

bound on the total running time, no more than 10 minutes was actually required for our method to converge in all those instances. As said before, this statistic has to be read with caution since the computing environments used were not comparable. A deeper look into the time performance of our method, however, dismisses the hasty conclusion that this difference might explain its higher success rate. In the histogram of Figure 11, we show the total number of instances solved by our algorithm within certain discrete amounts of time. Note the remarkable fact that $\sim 90\%$ of the instances were solved roughly within 2 minutes.

Table IV. Summary of Results

Source	Instance Group	n	Optimality Rate	Guards	Iterations	Time (s)
Couto et al.	Simple (30 inst. per size)	100	100.00%	13.17	1.93	0.94
		1,000	100.00%	126.57	4.50	181.27
		2,500	76.67%	312.39	5.78	2,349.84
	Orthogonal (30 inst. per size)	100	100.00%	14.63	2.50	0.64
		1,000	100.00%	147.90	5.40	120.29
		2,500	100.00%	371.40	6.87	1,139.80
	Von Koch (30 inst. per size)	100	100.00%	7.83	1.70	0.87
		1,000	100.00%	56.80	2.37	723.54
		1,500	73.33%	95.00	2.67	2,584.83
Bottino and Laurentini	Simple (20 inst. per size)	40	100.00%	5.55	1.15	0.10
		50	100.00%	6.60	1.60	0.24
		60	100.00%	8.35	1.30	0.27
	Orthogonal (20 inst. per size)	40	100.00%	6.00	1.40	0.07
		50	100.00%	7.70	1.55	0.12
		60	100.00%	9.10	1.50	0.16
Kröller et al.	Spikes (30 inst. per size)	100	100.00%	4.37	1.07	1.40
		200	100.00%	6.87	1.00	7.30
		500	100.00%	9.77	1.00	173.23
Tozoni et al.	Simple-simple (30 inst. per size)	100	100.00%	13.13	2.27	4.92
		200	100.00%	24.30	3.47	34.68
		500	93.33%	59.68	5.82	498.69
		1,000	40.00%	115.08	5.42	2,139.29
	Ortho-ortho (30 inst. per size)	100	100.00%	12.27	2.97	7.88
		200	100.00%	24.97	4.90	45.29
		500	93.33%	62.46	6.29	433.03
		1,000	63.33%	127.74	7.79	2,097.72

5.7. General Experiments

Having analyzed our method in comparison to other approaches, let us focus on the overall performance of our algorithm on instances from all available sources.

First, Table IV sums up the results for polygons of three or four different sizes from each instance group, in order to characterize the behavior of our method on the widest group of instances possible.

Clearly, some classes of polygons seem easier than others: orthogonal polygons tend to take less time and reach a higher optimality rate than simple or von Koch polygons. The latter are clearly the hardest ones: within the average time taken to solve a von Koch instance of 1,000 vertices, we were able to solve simple polygons of double that size. In the same vein, it is evident that polygons with holes are at least an order of magnitude harder than their hole-free counterparts due to the higher density of the arrangements induced by the holes.

Although average time information can be useful for a general analysis, it is often necessary to observe the behavior on particular instances. In Figure 12, three box-plot charts present the running time results for different classes of polygons. In these charts, it is possible to perceive that the average times (from Table IV) for a given class and size of polygons may be quite disparate from the median, or even larger than the third quartile, due to outliers. For instance, consider the Ortho-ortho polygons with 200 vertices: the average time to solve the AGP is 45.29 seconds, while the median and the third quartile are 24.96 and 41.38 seconds, respectively. It is clear from these numbers that the instance classes contain polygons that differ considerably from each other, leading to AGP instances with very diverse degrees of difficulty. The high success

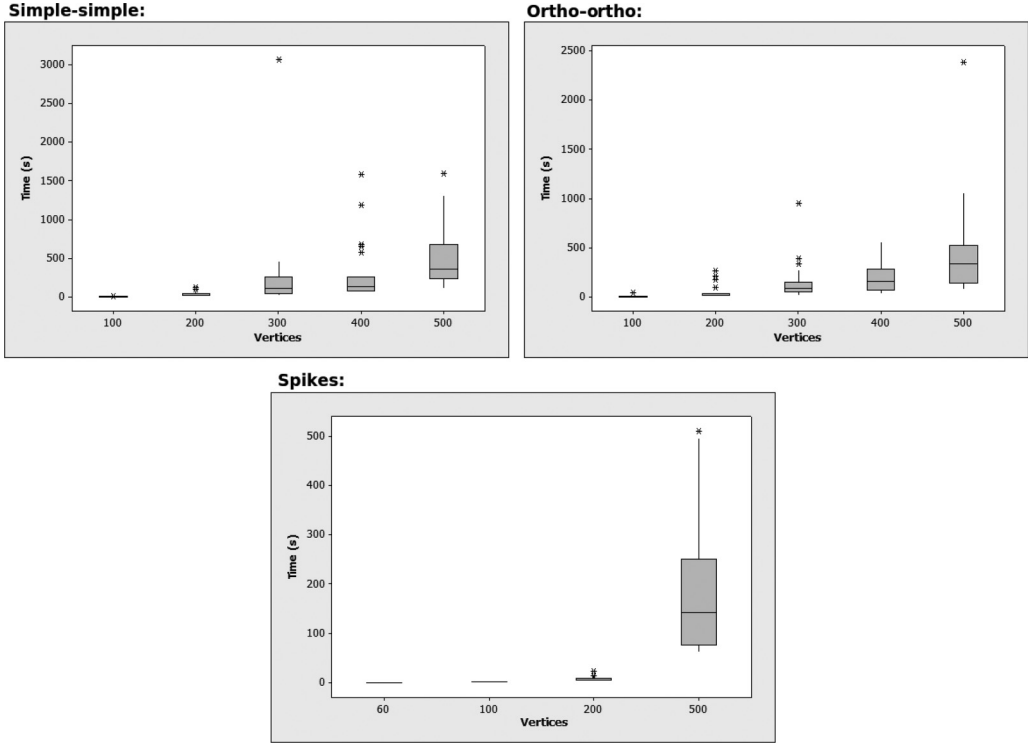


Fig. 12. Box-plot charts summarizing information about the running times to solve Simple-simple, Ortho-ortho, and Spike instances.

rate of our algorithm in such a heterogeneous benchmark indicates that the method is very robust.

Quite informative also is an analysis of how the total running time is divided between the different computational tasks, in practice. While one would expect that the larger amount of time should be spent solving the ILP (SCP) instances (after all, these are NP-hard problems), we observe from the histogram in Figure 13 that the time spent solving ILP models is usually lower than half of the total time. This can be credited to the high quality of modern solvers, such as XPRESS, but, in our case, also to the improvements made in our implementation, such as matrix reduction and the Lagrangian heuristic. Of course, were we to consider ever-growing instance sizes, this behavior is expected to change, asymptotically.

Needless to say, the amount of time spent in the resolution of ILP models is proportionally larger for instances with holes as a consequence of the denser and more complex arrangements, which led to larger and harder ILP instances.

5.8. Implementation

As mentioned earlier, our implementation relies on an ILP solver to quickly solve SCP instances. Initially, our choice was for the XPRESS Optimization Suite [XPRESS 2013] (version 7.0) as the state-of-the-art software capable of finding solutions for very large and complex instances in a matter of minutes. For this reason, all experiments reported up until now were conducted with XPRESS.

However, as XPRESS is a proprietary software, we also included in our implementation the possibility of using a free ILP solver to facilitate further tests and comparisons of

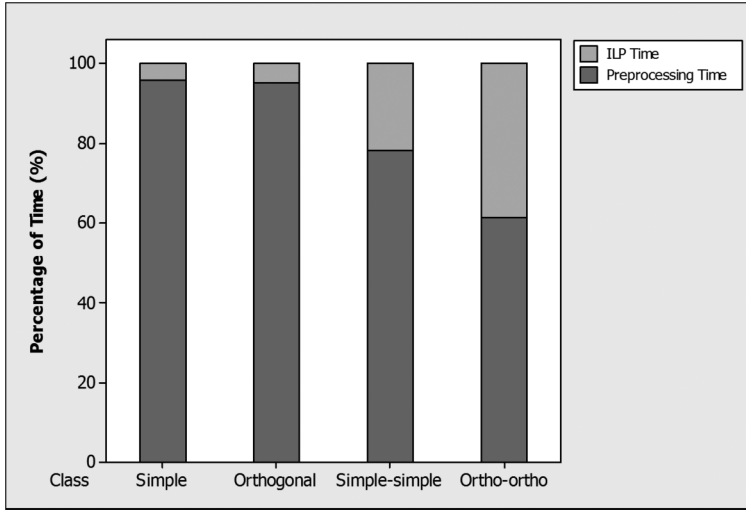


Fig. 13. Comparison between preprocessing time and solver time for polygons of 500 vertices.

Table V. Rates Between the Running Times When Employing XPRESS and GLPK

ILP Solver	Time Rate	
	Simple	Orthogonal
GLPK/XPRESS	2.56	0.89

our techniques with new algorithms. The free package of choice was the GLPK (GNU Linear Programming Kit) [GLPK 2013], a set of routines written in the ANSI C programming language and organized in the form of a callable library. Although switching to a noncommercial solver is expected to lead to some degradation in performance, it remains a handy and easy-to-use means for benchmarking future research on the subject.

As an illustration of the loss of performance, we carried out some extra tests to assess the impact of switching between these two ILP solvers. Table V shows the ratio between the time needed to solve instances of polygons with 1,000 vertices when running our code with XPRESS and with GLPK. From these results, we can see a considerable advantage when selecting XPRESS, at least for simple polygons. Curiously, when we consider orthogonal polygons, GLPK presents a modest advantage.

Although the total time ratio is a plausible way to measure how the choice of the ILP solver affects the program's performance, another conceivable evaluation is to calculate the fraction of total time spent on actually solving ILPs. This can be appraised from Figure 14, which shows the percentage of time used for computing SCP instances of simple and orthogonal polygons of 1,000 vertices. The dark and light gray bars represent the percentage of time spent on preprocessing and on ILP computation, respectively, when employing GLPK. As a comparative reference, the dashed lines indicate where the boundaries between the two bars are when XPRESS is applied. In view of what was reported in Section 5.7, one can see that the ILP computation has a greater influence on the total time when GLPK is used.

Moreover, a considerable loss of performance was observed in preliminary experiments conducted with polygons with holes when using GLPK.

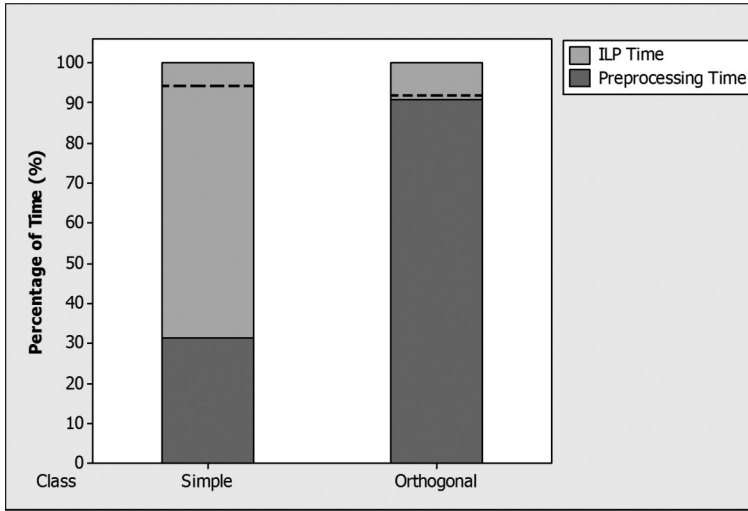


Fig. 14. Percentage of time spent in the preprocessing and ILP solver phases, for polygons of 1,000 vertices, when using GLPK. The dashed line indicates the corresponding break point for XPRSS.

6. CONCLUSION

For many years, much theoretical knowledge has been gathered on a number of hard geometric problems. More recently, the need for solving these problems to optimality in practice has gained much interest. In particular, the use of ILP formulations associated with state-of-the-art solvers has proven very successful in attaining exact solutions for large instances efficiently.

Following the trend of Couto et al. [2011], Kröller et al. [2012], and Tozoni et al. [2013], this article shows that this approach is proficuous to solving the Art Gallery Problem with point guards to optimality. The ILP model presented here has been implemented and tested in a total of 2,740 instances, and for all but 120 of them, optimal solutions were obtained in less than an hour of computational time.

While convergence remains dependent on the initial discretization, the current results already show an efficacy rate of more than 95%, which is an achievement unattained up to this point in time.

It remains a challenge for the future to study additional geometric properties that may assist in designing a discretization strategy that leads to proven convergence of the method described in this work. Answering the following questions might shed some light on the issue of convergence: Does there always exist an optimal rational solution whenever the vertices of the input polygon have only rational coordinates? Moreover, how can this help determine an ideal initial discretization?

Along with this article, we are making the code of our program available to interested researchers and practitioners. To facilitate its use, we implemented an option for selecting GLPK (GNU Linear Programming Kit) [GLPK 2013] as the ILP solver used for solving SCPs. As is the case of the required CGAL [CGAL 2012] library, GLPK also provides a free license. Therefore, to anyone willing to use our software, it will suffice to download those packages and compile the code on the preferred platform.

Although some loss of performance is perceived upon replacement of the commercial solver by GLPK, our implementation remains an invaluable tool for those interested in studying the AGP.

To the best of our knowledge, no other robust code that tackles this problem has ever been made accessible. Therefore, our aim in doing so is to provide a comparative gauge for reliable evaluation of different solution methods for the AGP.

ACKNOWLEDGMENTS

We would like to thank Alexander Kröller and Sandór Fekete for fruitful discussions and for providing additional sets of instances for our experiments. We are also grateful to Andrea Bottino for promptly dispensing the instances used in his work for the tests conducted in our investigation.

REFERENCES

- Yoav Amit, Joseph S. B. Mitchell, and Eli Packer. 2007. Locating guards for visibility coverage of polygons. In *Proceedings of the 9th Workshop on Algorithm Engineering and Experiments (ALENEX'07)*. SIAM, New Orleans, LA, 1–15.
- John E. Beasley. 1993. Lagrangian relaxation. In *Modern Heuristic Techniques for Combinatorial Problems*, Colin R. Reeves (Ed.). John Wiley & Sons, New York, NY, 243–303. <http://dl.acm.org/citation.cfm?id=166648.166660>
- Dorit Borrmann, Pedro J. de Rezende, Cid C. de Souza, Sándor P. Fekete, Stephan Friedrichs, Alexander Kröller, Andreas Nüchter, Christiane Schmidt, and Davi C. Tozoni. 2013. Point guards and point clouds: Solving general art gallery problems. In *Proceedings of the 29th Annual Symposium on Computational Geometry (SoCG'13)*. ACM, New York, NY, 347–348. DOI: <http://dx.doi.org/10.1145/2462356.2462361>
- Prosenjit Bose, Anna Lubiw, and J. Ian Munro. 2002. Efficient visibility queries in simple polygons. *Computational Geometry* 23, 3 (2002), 313–335.
- Andrea Bottino and Aldo Laurentini. 2011. A nearly optimal algorithm for covering the interior of an art gallery. *Pattern Recognition* 44, 5 (2011), 1048–1056. DOI: <http://dx.doi.org/10.1016/j.patcog.2010.11.010>
- CGAL. 2012. Computational Geometry Algorithms Library. (2012). www.cgal.org (last accessed January 2012).
- Kyung-Yong Chwa, Byung-Cheol Jo, Christian Knauer, Esther Moet, René van Oostrum, and Chan-Su Shin. 2006. Guarding art galleries by guarding witnesses. *International Journal of Computational Geometry And Applications* 16, 02n03 (2006), 205–226. DOI: <http://dx.doi.org/10.1142/S0218195906002002>
- Marcelo C. Couto, Pedro J. de Rezende, and Cid C. de Souza. 2011. An exact algorithm for minimizing vertex guards on art galleries. *International Transactions in Operational Research* 18, 4 (2011), 425–448. DOI: <http://dx.doi.org/10.1111/j.1475-3995.2011.00804.x>
- Marcelo C. Couto, Davi C. Tozoni, Pedro J. de Rezende, and Cid C. de Souza. 2013. The Art Gallery Problem Project (point guards). (2013). www.ic.unicamp.br/~cid/Problem-instances/Art-Gallery.
- Sándor P. Fekete, Stephan Friedrichs, Alexander Kröller, and Christiane Schmidt. 2015. Facets for art gallery problems. *Algorithmica* 73, 2 (2015), 411–440. DOI: <http://dx.doi.org/10.1007/s00453-014-9961-x>
- Subir K. Ghosh. 1987. Approximation algorithms for art gallery problems. In *Proceedings of the Canadian Information Processing Society Congress*. Canadian Information Processing Society, Mississauga, Ontario, Canada, 429–434.
- Subir K. Ghosh. 2007. *Visibility Algorithms in the Plane*. Cambridge University Press, New York, NY.
- Subir K. Ghosh. 2010. Approximation algorithms for art gallery problems in polygons. *Discrete Applied Mathematics* 158, 6 (2010), 718–722. DOI: <http://dx.doi.org/10.1016/j.dam.2009.12.004>
- GLPK. 2013. *GNU Linear Programming Kit*. GNU. <http://www.gnu.org/software/glpk/> (access December 2013).
- Alexander Kröller, Tobias Baumgartner, Sándor P. Fekete, and Christiane Schmidt. 2012. Exact solutions and bounds for general art gallery problems. *Journal on Experimental Algorithmics* 17, 1, Article 2.3 (May 2012), 1.13 pages. DOI: <http://dx.doi.org/10.1145/2133803.2184449>
- D. T. Lee and Arthur Lin. 1986. Computational complexity of art gallery problems. *IEEE Transactions on Information Theory*, 32, 2 (March 1986), 276–282. DOI: <http://dx.doi.org/10.1109/TIT.1986.1057165>
- Joseph O'Rourke. 1987. *Art Gallery Theorems and Algorithms*. Oxford University Press, New York, NY.
- Thomas C. Shermer. 1992. Recent results in art galleries. *Proceedings of the IEEE* 80, 9 (September 1992), 1384–1399. DOI: <http://dx.doi.org/10.1109/5.163407>
- Davi C. Tozoni, Pedro J. de Rezende, and Cid C. de Souza. 2013. The quest for optimal solutions for the art gallery problem: A practical iterative algorithm. In *Proceedings of the 12th International Symposium on Experimental Algorithms (SEA'13) (Lecture Notes in Computer Science)*, Vincenzo Bonifaci, Camil Demetrescu, and Alberto Marchetti-Spaccamela (Eds.), Vol. 7933. Springer, Rome, Italy, 320–336.

- Jorge Urrutia. 2000. Art gallery and illumination problems. In *Handbook of Computational Geometry*, J. R. Sack and J. Urrutia (Eds.). Elsevier Science Publishers, Amsterdam, 973–1027.
- Jan van Leeuwen and Anneke A. Schoone. 1981. Untangling a traveling salesman tour in the plane. In *Proceedings of the 7th Conference on Graph-Theoretic Concepts in Computer Science*. Rijksuniversiteit Vakgroep Informatica, Munich, 87–98. <http://books.google.com.br/books?id=LsEuHAAACAAJ>.
- XPRESS. 2013. *Xpress Optimization Suite*. FICO Solutions. <http://www.fico.com/en/products/fico-xpress-optimization-suite> (access November 2013).

Received January 2015; revised January 2016; accepted February 2016